

Week 08 Lectures

Assignment 2

1/141

Aim: experimental analysis of signature-based indexing

- tuple-level superimposed codeword signatures
- page-level superimposed codeword signatures
- bit-sliced superimposed codeword signatures

Large numbers of tuples, inserted into a *relation*:

- implemented as **one data file** plus **three signature files**

Produce several instance of (data+signatures) relations

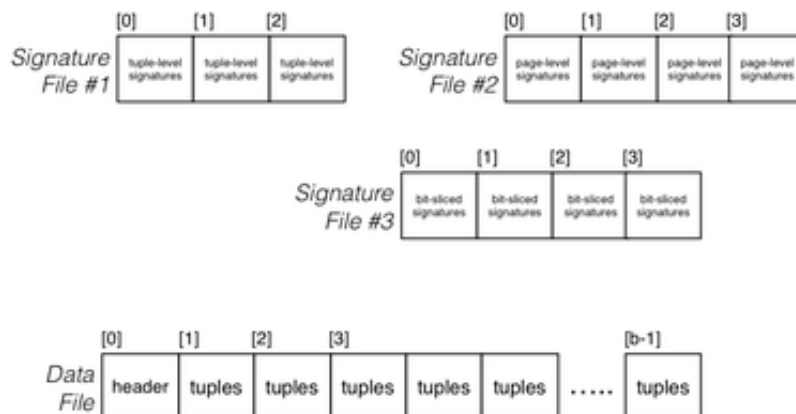
Run a set of benchmark PMR queries on these relations

Measure query costs and analyse

... Assignment 2

2/141

File structures:



... Assignment 2

3/141

We supply:

- a program to generate pseudo-random tuples (lots of them)
 - where tuples look like: (id, name, *numeric-attributes*)
- a hash function (from PostgreSQL)

You write:

- a program to build (data+signature) files from tuples
- a program to run queries (val,?,?,val,?,...) against the data

Your programs must be parameterized:

- builder accepts different values for signature parameters
- query processor specifies what kind of signatures to use

... Assignment 2

4/141

Experiments:

- run each query using each signature type
- for each sig type
 - read appropriate signature pages (maybe all)
 - determine which data pages to read
 - fetch pages and search for matching tuples
 - determine total cost (#signature pages + #data pages)
- record costs and analyse/compare

Query Processing So Far

5/141

Steps in processing an SQL statement

- parse, map to relation algebra (RA) expression
- transform to more efficient RA expression
- instantiate RA operators to DBMS operations
- execute DBMS operations (aka query plan)

Cost-based optimisation:

- generate possible query plans (via rewriting/heuristics)
- estimate cost of each plan (sum costs of operations)
- choose the lowest-cost plan (... and choose quickly)

Expression Rewriting Rules

6/141

Since RA is a well-defined formal system

- there exist many algebraic laws on RA expressions
- which can be used as a basis for expression rewriting
- in order to produce *equivalent* (more-efficient?) expressions

Expression transformation based on such rules can be used

- to simplify/improve SQL \rightarrow RA mapping results
- to generate new plan variations to check in query optimisation

Relational Algebra Laws

7/141

Commutative and Associative Laws:

- $R \bowtie S \leftrightarrow S \bowtie R$, $(R \bowtie S) \bowtie T \leftrightarrow R \bowtie (S \bowtie T)$ (natural join)
- $R \cup S \leftrightarrow S \cup R$, $(R \cup S) \cup T \leftrightarrow R \cup (S \cup T)$
- $R \bowtie_{Cond} S \leftrightarrow S \bowtie_{Cond} R$ (theta join)
- $\sigma_c(\sigma_d(R)) \leftrightarrow \sigma_d(\sigma_c(R))$

Selection splitting (where c and d are conditions):

- $\sigma_{c \wedge d}(R) \leftrightarrow \sigma_c(\sigma_d(R))$
- $\sigma_{c \vee d}(R) \leftrightarrow \sigma_c(R) \cup \sigma_d(R)$

... Relational Algebra Laws

8/141

Selection pushing ($\sigma_c(R \cup S)$ and $\sigma_c(R \cap S)$):

- $\sigma_c(R \cup S) \leftrightarrow \sigma_c R \cup \sigma_c S$, $\sigma_c(R \cap S) \leftrightarrow \sigma_c R \cap \sigma_c S$

Selection pushing with join ...

- $\sigma_c(R \bowtie S) \leftrightarrow \sigma_c(R) \bowtie S$ (if c refers only to attributes from R)
- $\sigma_c(R \bowtie S) \leftrightarrow R \bowtie \sigma_c(S)$ (if c refers only to attributes from S)

If c contains attributes from both R and S :

- $\sigma_{c' \wedge c''}(R \bowtie S) \leftrightarrow \sigma_{c'}(R) \bowtie \sigma_{c''}(S)$
- c' contains only R attributes, c'' contains only S attributes

... Relational Algebra Laws

9/141

Rewrite rules for projection ...

All but last projection can be ignored:

- $\pi_{L1}(\pi_{L2}(\dots \pi_{Ln}(R))) \rightarrow \pi_{L1}(R)$

Projections can be pushed into joins:

- $\pi_L(R \bowtie_c S) \leftrightarrow \pi_L(\pi_M(R) \bowtie_c \pi_N(S))$

where

- M and N must contain all attributes needed for c
- M and N must contain all attributes used in L ($L \subset M \cup N$)

... Relational Algebra Laws

10/141

Subqueries \Rightarrow convert to a join

Example: (on schema Courses(id,code,...), Enrolments(cid,sid,...), Students(id,name,...))

```
select c.code, count(*)
from   Courses c
where  c.id in (select cid from Enrolments)
group  by c.code
```

becomes

```
select c.code, count(*)
from   Courses c join Enrolments e on c.id = e.cid
group  by c.code
```

but not

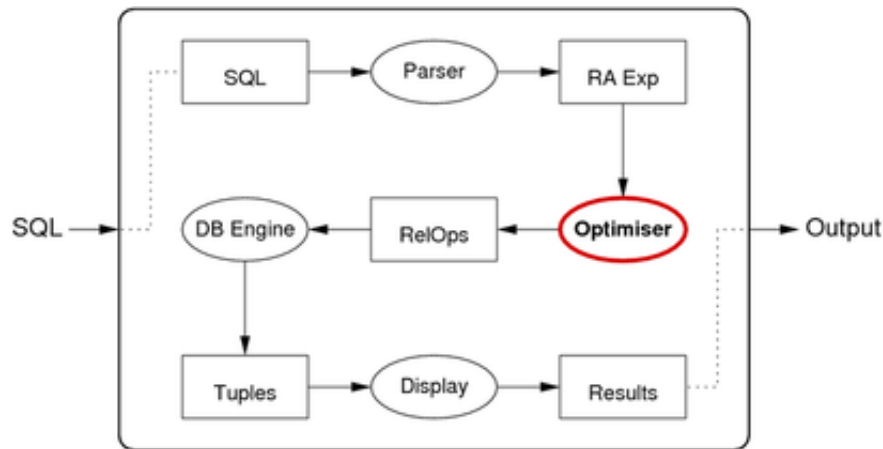
```
select e.sid as student_id, e.cid as course_id
from   Enrolments e
where  e.sid = (select max(id) from Students)
```

Query Optimisation

Query Optimisation

12/141

Query optimiser: RA expression → efficient evaluation plan



... Query Optimisation

13/141

Query optimisation is a critical step in query evaluation.

The query optimiser

- takes relational algebra expression from SQL compiler
- produces sequence of RelOps to evaluate the expression
- *query execution plan* should provide efficient evaluation

"Optimisation" is a misnomer since query optimisers

- aim to find a good plan ... but maybe not optimal

Observed Query Time = Planning time + Evaluation time

... Query Optimisation

14/141

Why do we not generate optimal query execution plans?

Finding an optimal query plan ...

- requires exhaustive search of a *space of possible plans*
- for each possible plan, need to estimate cost (not cheap)

Even for relatively small query, search space is *very large*.

Compromise:

- do limited search of query plan space (guided by heuristics)
- *quickly* choose a *reasonably efficient* execution plan

15/141

Approaches to Optimisation

Three main classes of techniques developed:

- algebraic (equivalences, rewriting, heuristics)
- physical (execution costs, search-based)
- semantic (application properties, heuristics)

All driven by aim of minimising (or at least reducing) "cost".


Real query optimisers use a combination of algebraic+physical.

Semantic QO is good idea, but expensive/difficult to implement.

... Approaches to Optimisation

16/141

Example of optimisation transformations:



For join, may also consider sort/merge join and hash join.

Cost-based Query Optimiser

17/141

Approximate algorithm for cost-based optimisation:

```

translate SQL query to RAexp
for enough transformations RA' of RAexp {
  while (more choices for RelOps) {
    plan = {}; i = 0
    for each node e of RA' (recursively) {
      select RelOp method for e
      plan[i++] = RelOp method for e
    }
    cost = 0
    for each op in plan[] { cost += Cost(op) }
    if (cost < MinCost)
      { MinCost = cost; BestPlan = plan }
  }
}

```

Heuristics: push selections down, consider only left-deep join trees.

Exercise 1: Alternative Join Plans

18/141

Consider the schema

```
Students(id,name,...)   Enrol(student,course,mark)
Staff(id,name,...)      Courses(id,code,term,lic,...)
```

the following query on this schema

```
select c.code, s.id, s.name
from   Students s, Enrol e, Courses c, Staff f
where  s.id=e.student and e.course=c.id
      and c.lic=f.id and c.term='11s2'
      and f.name='John Shepherd'
```

Show some possible evaluation orders for this query.

Cost Models and Analysis

19/141

The cost of evaluating a query is determined by:

- size of relations (database relations and temporary relations)
- access mechanisms (indexing, hashing, sorting, join algorithms)
- size/number of main memory buffers (and replacement strategy)

Analysis of costs involves *estimating*:

- size of intermediate results
- number of secondary storage accesses

Choosing Access Methods (RelOps)

20/141

Performed for each node in RA expression tree ...

Inputs:

- a single RA operation (σ , π , \bowtie)
- information about file organisation, data distribution, ...
- list of operations available in the database engine

Output:

- specific DBMS operation to implement this RA operation

... Choosing Access Methods (RelOps)

21/141

Example:

- RA operation: $Sel_{[name='John' \wedge age>21]}(Student)$
- Student relation has B-tree index on name
- database engine (obviously) has B-tree search method

giving

```
tmp[i]   := BtreeSearch[name='John'](Student)
tmp[i+1] := LinearSearch[age>21](tmp[i])
```

Where possible, use pipelining to avoid storing `tmp[i]` on disk.

... Choosing Access Methods (RelOps)

22/141

Rules for choosing σ access methods:

- $\sigma_{A=c}(R)$ and R has index on $A \Rightarrow \text{indexSearch}[A=c](R)$
- $\sigma_{A=c}(R)$ and R is hashed on $A \Rightarrow \text{hashSearch}[A=c](R)$
- $\sigma_{A=c}(R)$ and R is sorted on $A \Rightarrow \text{binarySearch}[A=c](R)$
- $\sigma_{A \geq c}(R)$ and R has clustered index on A
 $\Rightarrow \text{indexSearch}[A=c](R)$ then scan
- $\sigma_{A \geq c}(R)$ and R is hashed on A
 $\Rightarrow \text{linearSearch}[A \geq c](R)$

... Choosing Access Methods (RelOps)

23/141

Rules for choosing \bowtie access methods:

- $R \bowtie S$ and R fits in memory buffers $\Rightarrow \text{bnlJoin}(R, S)$
- $R \bowtie S$ and S fits in memory buffers $\Rightarrow \text{bnlJoin}(S, R)$
- $R \bowtie S$ and R, S sorted on join attr $\Rightarrow \text{smJoin}(R, S)$
- $R \bowtie S$ and R has index on join attr $\Rightarrow \text{inlJoin}(S, R)$
- $R \bowtie S$ and no indexes, no sorting $\Rightarrow \text{hashJoin}(R, S)$

(bnl = block nested loop; inl = index nested loop; sm = sort merge)

Cost Estimation

24/141

Without executing a plan, cannot always know its precise cost.

Thus, query optimisers *estimate* costs via:

- cost of performing operation (dealt with in earlier lectures)
- size of result (which affects cost of performing next operation)

Result size determined by statistical measures on relations, e.g.

r_S	cardinality of relation S
R_S	avg size of tuple in relation S
$V(A, S)$	# distinct values of attribute A
$\min(A, S)$	min value of attribute A
$\max(A, S)$	max value of attribute A

Estimating Projection Result Size

25/141

Straightforward, since we know:

- number of tuples in output

$$r_{out} = |\pi_{a,b,\dots}(T)| = |T| = r_T \quad (\text{in SQL, because of bag semantics})$$

- size of tuples in output

$$R_{out} = \text{sizeof}(a) + \text{sizeof}(b) + \dots + \text{tuple-overhead}$$

Assume page size B , $b_{out} = \lceil r_T / c_{out} \rceil$, where $c_{out} = \lfloor B / R_{out} \rfloor$

If using `select distinct ...`

- $|\pi_{a,b,\dots}(T)|$ depends on proportion of duplicates produced

Estimating Selection Result Size

26/141

Selectivity = fraction of tuples expected to satisfy a condition.

Common assumption: attribute values uniformly distributed.

Example: Consider the query

```
select * from Parts where colour='Red'
```

If $V(\text{colour}, \text{Parts})=4$, $r=1000 \Rightarrow |\sigma_{\text{colour}=\text{red}}(\text{Parts})|=250$

In general, $|\sigma_{A=c}(R)| \approx r_R / V(A, R)$

Heuristic used by PostgreSQL: $|\sigma_{A=c}(R)| \approx r/10$

... Estimating Selection Result Size

27/141

Estimating size of result for e.g.

```
select * from Enrolment where year > 2005;
```

Could estimate by using:

- uniform distribution assumption, r , min/max years

Assume: $\min(\text{year})=2000$, $\max(\text{year})=2009$, $|\text{Enrolment}|=10^5$

- 10^5 from 2000-2009 means approx 10000 enrolments/year
- this suggests 40000 enrolments since 2006

Heuristic used by some systems: $|\sigma_{A>c}(R)| \approx r/3$

... Estimating Selection Result Size

28/141

Estimating size of result for e.g.

```
select * from Enrolment where course <> 'COMP9315';
```

Could estimate by using:

- uniform distribution assumption, r , domain size

e.g. $V(\text{course}, \text{Enrolment}) = 2000$, $|\sigma_{A \neq c}(E)| = r * 1999/2000$

Heuristic used by some systems: $| \sigma_{A \prec c}(R) | \approx r$

Exercise 2: Selection Size Estimation

29/141

Assuming that

- all attributes have uniform distribution of data values
- attributes are independent of each other

Give formulae for the number of expected results for

1. `select * from R where not A=k`
2. `select * from R where A=k and B=j`
3. `select * from R where A in (k,l,m,n)`

where j, k, l, m, n are constants.

Assume: $V(A,R) = 10$ and $V(B,R)=100$ and $r=1000$

... Estimating Selection Result Size

30/141

How to handle non-uniform attribute value distributions?

- collect statistics about the values stored in the attribute/relation
- store these as e.g. a histogram in the meta-data for the relation

So, for part colour example, might have distribution like:

White: 35% Red: 30% Blue: 25% Silver: 10%

Use histogram as basis for determining # selected tuples.

Disadvantage: cost of storing/maintaining histograms.

... Estimating Selection Result Size

31/141

Summary: analysis relies on operation and data distribution:

E.g. `select * from R where a = k;`

Case 1: $uniq(R.a) \Rightarrow 0 \text{ or } 1 \text{ result}$

Case 2: $r_R \text{ tuples \&\& } size(dom(R.a)) = n \Rightarrow r_R / n \text{ results}$

E.g. `select * from R where a < k;`

Case 1: $k \leq min(R.a) \Rightarrow 0 \text{ results}$

Case 2: $k > max(R.a) \Rightarrow \approx r_R \text{ results}$

Case 3: $size(dom(R.a)) = n \Rightarrow ? min(R.a) \dots k \dots max(R.a) ?$

Estimating Join Result Size

32/141

Analysis relies on semantic knowledge about data/relations.

Consider equijoin on common attr: $R \bowtie_a S$

Case 1: $values(R.a) \cap values(S.a) = \{\} \Rightarrow size(R \bowtie_a S) = 0$

Case 2: $uniq(R.a) \text{ and } uniq(S.a) \Rightarrow size(R \bowtie_a S) \leq \min(|R|, |S|)$

Case 3: $pkey(R.a) \text{ and } fkey(S.a) \Rightarrow size(R \bowtie_a S) \leq |S|$

Exercise 3: Join Size Estimation

33/141

How many tuples are in the output from:

1. `select * from R, S where R.s = S.id`
where `S.id` is a primary key and `R.s` is a foreign key referencing `S.id`
2. `select * from R, S where R.s <> S.id`
where `S.id` is a primary key and `R.s` is a foreign key referencing `S.id`
3. `select * from R, S where R.x = S.y`
where `R.x` and `S.y` have no connection except that $dom(R.x) = dom(S.y)$

Under what conditions will the first query have maximum size?

Cost Estimation: Postscript

34/141

Inaccurate cost estimation can lead to poor evaluation plans.

Above methods can (sometimes) give inaccurate estimates.

To get more accurate cost estimates:

- more time ... complex computation of selectivity
- more space ... storage for histograms of data values

Either way, optimisation process costs more (more than query?)

Trade-off between optimiser performance and query performance.

PostgreSQL Query Optimiser

Overview of QOpt Process

36/141

Input: tree of **Query** nodes returned by parser

Output: tree of **Plan** nodes used by query *executor*

- wrapped in a **PlannedStmt** node containing state info

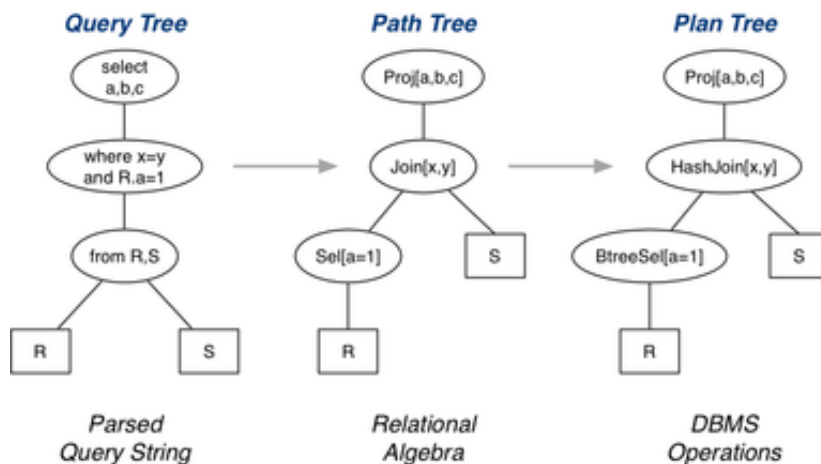
Intermediate data structures are trees of **Path** nodes

- a path tree represents one evaluation order for a query

All **Node** types are defined in `include/nodes/*.h`

... Overview of QOpt Process

37/141



QOpt Data Structures

38/141

Generic **Path** node structure:

```
typedef struct Path
{
    NodeTag    type;           // scan/join/...
    NodeTag    pathtype;       // specific method
    RelOptInfo *parent;        // output relation
    PathTarget *pathtarget;    // list of Vars/Exprs, width
    // estimated execution costs for path ...
    Cost       startup_cost;   // setup cost
    Cost       total_cost;     // total cost
    List       *pathkeys;      // sort order
} Path;
```

PathKey = (opfamily:Oid, strategy:SortDir, nulls_first:bool)

... QOpt Data Structures

39/141

Specialised **Path** nodes (simplified):

```
typedef struct IndexPath
{
    Path    path;
    IndexOptInfo *indexinfo; // index for scan
    List    *indexclauses;   // index select conditions
    ...
    ScanDirection indexscandir; // used by planner
    Selectivity indexselectivity; // estimated #results
} IndexPath;
```

```
typedef struct JoinPath
{
    Path    path;
    JoinType jointype;       // inner/outer/semi/anti
    Path    *outerpath;      // outer part of the join
    Path    *innerpath;      // inner part of the join
}
```

```
List      *restrictinfo; // join condition(s)
} JoinPath;
```

Query Optimisation Process

40/141

Query optimisation proceeds in two stages (after parsing)...

Rewriting:

- uses PostgreSQL's *rule* system
- query tree is expanded to include e.g. view definitions

Planning and optimisation:

- using cost-based analysis of generated paths
- via one of *two* different path generators
- chooses least-cost path from all those considered

Then produces a **Plan** tree from the selected path.

Top-down Trace of QOpt

41/141

Top-level of query execution: **backend/tcop/postgres.c**

```
exec_simple_query(const char *query_string)
{
    // lots of setting up ... including starting xact
    parsetree_list = pg_parse_query(query_string);
    foreach(parsetree, parsetree_list) {
        // Query optimisation
        querytree_list = pg_analyze_and_rewrite(parsetree,...);
        plantree_list = pg_plan_queries(querytree_list,...);
        // Query execution
        portal = CreatePortal(...plantree_list...);
        PortalRun(portal,...);
    }
    // lots of cleaning up ... including close xact
}
```

Assumes that we are dealing with multiple queries (i.e. SQL statements)

... Top-down Trace of QOpt

42/141

pg_analyze_and_rewrite()

- take a parse tree (from SQL parser)
- transforms Parse tree into Query tree (SQL → RA)
- applies rewriting rules (e.g. views)
- returns a list of Query trees

Code in: **backend/tcop/postgres.c**

... Top-down Trace of QOpt

43/141

pg_plan_queries()

- takes a list of parsed/re-written queries
- plans each one via `planner()`
 - which invokes `subquery_planner()` on each query
- returns a list of query plans

Code in: `backend/optimizer/plan/planner.c`

... Top-down Trace of QOpt

44/141

`subquery_planner()`

- performs algebraic transformations/simplifications, e.g.
 - simplifies conditions in **where** clauses
 - converts sub-queries in **where** to top-level join
 - moves **having** clauses with no aggregate into **where**
 - flattens sub-queries in join list
 - simplifies join tree (e.g. removes redundant terms), etc.
- sets up canonical version of query for plan generation
- invokes `grouping_planner()` to produce best path

Code in: `backend/optimizer/plan/planner.c`

... Top-down Trace of QOpt

45/141

`grouping_planner()` produces plan for one SQL statement

- preprocesses target list for INSERT/UPDATE
- handles "planning" for extended-RA SQL constructs:
 - set operations: UNION/INTERSECT/EXCEPT
 - GROUP BY, HAVING, aggregations
 - ORDER BY, DISTINCT, LIMIT
- invokes `query_planner()` for select/join trees

Code in: `backend/optimizer/plan/planmain.c`

... Top-down Trace of QOpt

46/141

`query_planner()` produces plan for a select/join tree

- make list of tables used in query
- split **where** qualifiers ("quals") into
 - restrictions (e.g. `r.a=1`) ... for selections
 - joins (e.g. `s.id=r.s`) ... for joins
- search for quals to enable merge/hash joins
- invoke `make_one_rel()` to find best path/plan

Code in: `backend/optimizer/plan/planmain.c`

... Top-down Trace of QOpt

47/141

`make_one_rel()` generates possible plans, selects best

- generate scan and index paths for base tables
 - using of restrictions list generated above

- generate access paths for the entire join tree
 - recursive process, controlled by `make_rel_from_joinlist()`
- returns a single "relation", representing result set

Code in: `backend/optimizer/path/allpaths.c`

Join-tree Generation

48/141

`make_rel_from_joinlist()` arranges path generation

- switches between two possible path tree generators
- path tree generators finally return best cost path

Standard path tree generator (`standard_join_search()`):

- "exhaustively" generates join trees (like System R)
- starts with 2-way joins, finds best combination
- then adds extra table to give 3-table join, etc.

Code in: `backend/optimizer/path/{allpaths.c,joinrels.c}`

... Join-tree Generation

49/141

Genetic query optimiser (`geqo`):

- uses genetic algorithm (GA) to generate path trees
- handles joins involving > `geqo_threshold` relations
- goals of this approach:
 - find near-optimal solution
 - examine far less than entire search space

Code in: `backend/optimizer/geqo/*.c`

... Join-tree Generation

50/141

Basic idea of genetic algorithm:

```
Optimize(join)
{
    t = 0
    p = initialState(join) // pool of (random) join orders
    for (t = 0; t < #generations; t++) {
        p' = recombination(p) // get parts of good join orders
        p'' = mutation(p')    // generate new variations
        p = selection(p'', p)  // choose best join orders
    }
}
```

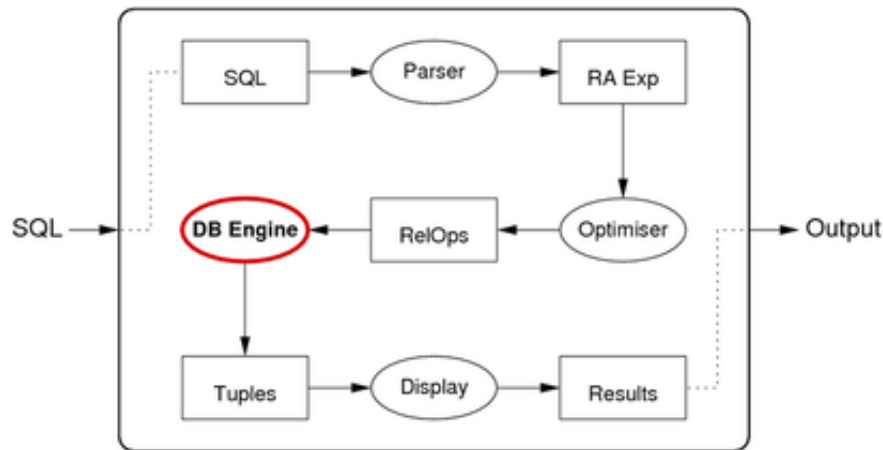
#generations determined by size of initial pool of join orders

Query Execution

Query Execution

52/141

Query execution: applies evaluation plan → result tuples



... Query Execution

53/141

Example of query translation:

```

select s.name, s.id, e.course, e.mark
from Student s, Enrolment e
where e.student = s.id and e.semester = '05s2';

```

maps to

$$\pi_{name,id,course,mark}(Stu \bowtie_{e.student=s.id} (\sigma_{semester=05s2} Enr))$$

maps to

```

Temp1 = BtreeSelect[semester=05s2](Enr)
Temp2 = HashJoin[e.student=s.id](Stu,Temp1)
Result = Project[name,id,course,mark](Temp2)

```

... Query Execution

54/141

A query execution plan:

- consists of a *collection of RelOps*
- executing together to produce a set of result tuples

Results may be passed from one operator to the next:

- *materialization* ... writing results to disk and reading them back
- *pipelining* ... generating and passing via memory buffers

Materialization

55/141

Steps in *materialization* between two operators

- first operator reads input(s) and writes results to disk
- next operator treats tuples on disk as its input
- in essence, the Temp tables are produced as real tables

Advantage:

- intermediate results can be placed in a file structure

(which can be chosen to speed up execution of subsequent operators)

Disadvantage:

- requires disk space/writes for intermediate results
- requires disk access to read intermediate results

Pipelining

56/141

How *pipelining* is organised between two operators:

- blocks execute "concurrently" as producer/consumer pairs
- first operator acts as producer; second as consumer
- structured as interacting iterators (open; while(next); close)

Advantage:

- no requirement for disk access (results passed via memory buffers)

Disadvantage:

- higher-level operators access inputs via linear scan, or
- requires sufficient memory buffers to hold all outputs

Iterators (reminder)

57/141

Iterators provide a "stream" of results:

- **iter = startScan(params)**
 - set up data structures for iterator (create state, open files, ...)
 - *params* are specific to operator (e.g. reln, condition, #buffers, ...)
- **tuple = nextTuple(iter)**
 - get the next tuple in the iteration; return null if no more
- **endScan(iter)**
 - clean up data structures for iterator

Other possible operations: reset to specific point, restart, ...

Pipelining Example

58/141


Consider the query:

```
select s.id, e.course, e.mark
from   Student s, Enrolment e
where  e.student = s.id and
       e.semester = '05s2' and s.name = 'John';
```

which maps to the RA expression

$$Proj_{[id, course, mark]}(Join_{[student=id]}(Sel_{[05s2]}(Enr), Sel_{[John]}(Stu)))$$


which could be represented by the RA expression tree



... Pipelining Example

59/141

Modelled as communication between RA tree nodes:



Note: likely that projection is combined with join in real DBMSs.

Disk Accesses

60/141

Pipelining cannot avoid all disk accesses.

Some operations use multiple passes (e.g. merge-sort, hash-join).

- data is written by one pass, read by subsequent passes

Thus ...

- *within* an operation, disk reads/writes are possible
- *between* operations, no disk reads/writes are needed

PostgreSQL Query Execution

PostgreSQL Query Execution

62/141

Defs: **src/include/executor** and **src/include/nodes**

Code: **src/backend/executor**

PostgreSQL uses pipelining ...

- query plan is a tree of **Plan** nodes
- each type of node implements one kind of RA operation (node implements specific access method via iterator interface)
- node types e.g. **Scan**, **Group**, **Indexscan**, **Sort**, **HashJoin**

- execution is managed via a tree of **PlanState** nodes
(mirrors the structure of the tree of Plan nodes; holds execution state)

PostgreSQL Executor

63/141

Modules in **src/backend/executor** fall into two groups:

execXXX (e.g. execMain, execProcnode, execScan)

- implement generic control of plan evaluation (execution)
- provide overall plan execution and dispatch to node iterators

nodeXXX (e.g. nodeSeqscan, nodeNestloop, nodeGroup)

- implement iterators for specific types of RA operators
- typically contains **ExecInitXXX**, **ExecXXX**, **ExecEndXXX**

... PostgreSQL Executor

64/141

Much simplified view of PostgreSQL executor:

```
ExecutePlan(execState, planStateNode, ...) {
    process "before each statement" triggers
    for (;;) {
        tuple = ExecProcNode(planStateNode)
        if (no more tuples) return END
        check tuple validity // MVCC
        if (got a tuple) break
    }
    process "after each statement" triggers
    return tuple
}
...
```

... PostgreSQL Executor

65/141

Executor overview (cont):

```
...
ExecProcNode(node) {
    switch (nodeType(node)) {
    case SeqScan:
        result = ExecSeqScan(node); break;
    case NestLoop:
        result = ExecNestLoop(node); break;
    ...
    }
    return result;
}
```

Example PostgreSQL Execution

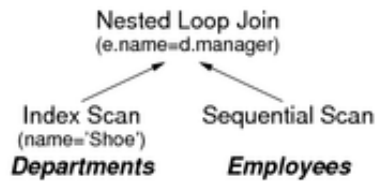
66/141

Consider the query:

```
-- get manager's age and # employees in Shoe department
select e.age, d.nemps
```

```
from Departments d, Employees e
where e.name = d.manager and d.name = 'Shoe'
```

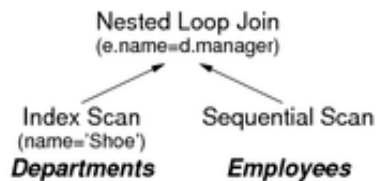
and its execution plan tree



... Example PostgreSQL Execution

67/141

The execution plan tree



contains three nodes:

- NestedLoop with join condition (Outer.manager = Inner.name)
- IndexScan on Departments with selection (name = 'Shoe')
- SeqScan on Employees

... Example PostgreSQL Execution

68/141

Initially `InitPlan()` invokes `ExecInitNode()` on plan tree root.

`ExecInitNode()` sees a `NestedLoop` node ...
 so dispatches to `ExecInitNestLoop()` to set up iterator
 then invokes `ExecInitNode()` on left and right sub-plans
 in left subPlan, `ExecInitNode()` sees an `IndexScan` node
 so dispatches to `ExecInitIndexScan()` to set up iterator
 in right sub-plan, `ExecInitNode()` sees a `SeqScan` node
 so dispatches to `ExecInitSeqScan()` to set up iterator

Result: a plan state tree with same structure as plan tree.

... Example PostgreSQL Execution

69/141

Execution: `ExecutePlan()` repeatedly invokes `ExecProcNode()`.

`ExecProcNode()` sees a `NestedLoop` node ...
 so dispatches to `ExecNestedLoop()` to get next tuple
 which invokes `ExecProcNode()` on its sub-plans
 in left sub-plan, `ExecProcNode()` sees an `IndexScan` node
 so dispatches to `ExecIndexScan()` to get next tuple
 if no more tuples, return END
 for this tuple, invoke `ExecProcNode()` on right sub-plan
`ExecProcNode()` sees a `SeqScan` node
 so dispatches to `ExecSeqScan()` to get next tuple
 check for match and return joined tuples if found
 continue scan until end
 reset right sub-plan iterator

Result: stream of result tuples returned via `ExecutePlan()`

Query Performance

Performance Tuning

71/141

How to make a database perform "better"?

Good performance may involve any/all of:

- making applications using the DB run faster
- lowering response time of queries/transactions
- improving overall transaction throughput

Remembering that, to some extent ...

- the query optimiser removes choices from DB developers
- by making its own decision on the optimal execution plan

... Performance Tuning

72/141

Tuning requires us to consider the following:

- which queries and transactions will be used?
(e.g. check balance for payment, display recent transaction history)
- how frequently does each query/transaction occur?
(e.g. 90% withdrawals; 10% deposits; 50% balance check)
- are there time constraints on queries/transactions?
(e.g. EFTPOS payments must be approved within 7 seconds)
- are there uniqueness constraints on any attributes?
(define indexes on attributes to speed up insertion uniqueness check)
- how frequently do updates occur?
(indexes slow down updates, because must update table *and* index)

... Performance Tuning

73/141

Performance can be considered at two times:

- *during* schema design
 - typically towards the end of schema design process
 - requires schema transformations such as *denormalisation*
- *outside* schema design
 - typically after application has been deployed/used
 - requires adding/modifying data structures such as *indexes*

Difficult to predict what query optimiser will do, so ...

- implement queries using methods which *should* be efficient
- observe execution behaviour and modify query accordingly

PostgreSQL Query Tuning

74/141

PostgreSQL provides the **explain** statement to

- give a representation of the query execution plan

- with information that may help to tune query performance

Usage:

`EXPLAIN [ANALYZE] Query`

Without ANALYZE, EXPLAIN shows plan with estimated costs.

With ANALYZE, EXPLAIN executes query and prints real costs.

Note that runtimes may show considerable variation due to buffering.

EXPLAIN Examples

75/141

Database

```
course_enrolments(student, course, mark, grade, ...)
courses(id, subject, semester, homepage)
people(id, family, given, title, name, ..., birthday)
program_enrolments(id, student, semester, program, wam, ...)
students(id, stype)
subjects(id, code, name, longname, uoc, offeredby, ...)
```

where

table_name	n_records
course_enrolments	525688
courses	73220
people	55767
program_enrolments	193456
students	31361
subjects	17779

... EXPLAIN Examples

76/141

Example: Select on non-indexed attribute

```
uni=# explain
uni=# select * from Students where stype='local';
          QUERY PLAN
-----
Seq Scan on students
    (cost=0.00..562.01 rows=23544 width=9)
    Filter: ((stype)::text = 'local'::text)
```

where

- Seq Scan = operation (plan node)
- cost=*StartupCost*..*TotalCost*
- rows=*NumberOfResultTuples*
- width=*SizeOfTuple* (# bytes)

... EXPLAIN Examples

77/141

More notes on explain output:

- each major entry corresponds to a plan node
 - e.g. Seq Scan, Index Scan, Hash Join, Merge Join, ...

- some nodes include additional qualifying information
 - e.g. Filter, Index Cond, Hash Cond, Buckets, ...
- cost values in explain are estimates (notional units)
- explain analyze also includes actual time costs (ms)
- costs of parent nodes include costs of all children
- estimates of #results based on sample of data

... EXPLAIN Examples

78/141

Example: Select on non-indexed attribute with actual costs

```
uni=# explain analyze
uni=# select * from Students where stype='local';
          QUERY PLAN
-----
Seq Scan on students
    (cost=0.00..562.01 rows=23544 width=9)
    (actual time=0.052..5.792 rows=23551 loops=1)
  Filter: ((stype)::text = 'local'::text)
  Rows Removed by Filter: 7810
 Planning time: 0.075 ms
 Execution time: 6.978 ms
```

... EXPLAIN Examples

79/141

Example: Select on indexed, unique attribute

```
uni=# explain analyze
uni=# select * from Students where id=100250;
          QUERY PLAN
-----
Index Scan using student_pkey on student
    (cost=0.00..8.27 rows=1 width=9)
    (actual time=0.049..0.049 rows=0 loops=1)
 Index Cond: (id = 100250)
 Total runtime: 0.1 ms
```

... EXPLAIN Examples

80/141

Example: Select on indexed, unique attribute

```
uni=# explain analyze
uni=# select * from Students where id=1216988;
          QUERY PLAN
-----
Index Scan using students_pkey on students
    (cost=0.29..8.30 rows=1 width=9)
    (actual time=0.011..0.012 rows=1 loops=1)
 Index Cond: (id = 1216988)
 Planning time: 0.066 ms
 Execution time: 0.026 ms
```

... EXPLAIN Examples

81/141

Example: Join on a primary key (indexed) attribute (2016)

```
uni=# explain analyze
uni=# select s.id,p.name
```

```

uni=# from Students s, People p where s.id=p.id;
      QUERY PLAN
-----
Hash Join (cost=988.58..3112.76 rows=31048 width=19)
  (actual time=11.504..39.478 rows=31048 loops=1)
  Hash Cond: (p.id = s.id)
    -> Seq Scan on people p
      (cost=0.00..989.97 rows=36497 width=19)
      (actual time=0.016..8.312 rows=36497 loops=1)
    -> Hash (cost=478.48..478.48 rows=31048 width=4)
      (actual time=10.532..10.532 rows=31048 loops=1)
      Buckets: 4096 Batches: 2 Memory Usage: 548kB
    -> Seq Scan on students s
      (cost=0.00..478.48 rows=31048 width=4)
      (actual time=0.005..4.630 rows=31048 loops=1)
Total runtime: 41.0 ms

```

... EXPLAIN Examples

82/141

Example: Join on a primary key (indexed) attribute (2018)

```

uni=# explain analyze
uni=# select s.id,p.name
uni=# from Students s, People p where s.id=p.id;
      QUERY PLAN
-----
Merge Join (cost=0.58..2829.25 rows=31361 width=18)
  (actual time=0.044..25.883 rows=31361 loops=1)
  Merge Cond: (s.id = p.id)
    -> Index Only Scan using students_pkey on students s
      (cost=0.29..995.70 rows=31361 width=4)
      (actual time=0.033..6.195 rows=31361 loops=1)
      Heap Fetches: 31361
    -> Index Scan using people_pkey on people p
      (cost=0.29..2434.49 rows=55767 width=18)
      (actual time=0.006..6.662 rows=31361 loops=1)
Planning time: 0.259 ms
Execution time: 27.327 ms

```

... EXPLAIN Examples

83/141

Example: Join on a non-indexed attribute (2016)

```

uni=# explain analyze
uni=# select s1.code, s2.code
uni=# from Subjects s1, Subjects s2
uni=# where s1.offeredBy=s2.offeredBy;
      QUERY PLAN
-----
Merge Join (cost=4449.13..121322.06 rows=7785262 width=18)
  (actual time=29.787..2377.707 rows=8039979 loops=1)
  Merge Cond: (s1.offeredby = s2.offeredby)
    -> Sort (cost=2224.57..2271.56 rows=18799 width=13)
      (actual time=14.251..18.703 rows=18570 loops=1)
      Sort Key: s1.offeredby
      Sort Method: external merge Disk: 472kB
    -> Seq Scan on subjects s1
      (cost=0.00..889.99 rows=18799 width=13)
      (actual time=0.005..4.542 rows=18799 loops=1)
    -> Sort (cost=2224.57..2271.56 rows=18799 width=13)
      (actual time=15.532..1100.396 rows=8039980 loops=1)
      Sort Key: s2.offeredby
      Sort Method: external sort Disk: 552kB
    -> Seq Scan on subjects s2

```

```

(cost=0.00..889.99 rows=18799 width=13)
(actual time=0.002..3.579 rows=18799 loops=1)
Total runtime: 2767.1 ms

```

... EXPLAIN Examples

84/141

Example: Join on a non-indexed attribute (2018)

```

uni=# explain analyze
uni=# select s1.code, s2.code
uni=# from Subjects s1, Subjects s2
uni=# where s1.offeredBy = s2.offeredBy;
               QUERY PLAN
-----
Hash Join  (cost=1286.03..108351.87 rows=7113299 width=18)
    (actual time=8.966..903.441 rows=7328594 loops=1)
    Hash Cond: (s1.offeredby = s2.offeredby)
    -> Seq Scan on subjects s1
        (cost=0.00..1063.79 rows=17779 width=13)
        (actual time=0.013..2.861 rows=17779 loops=1)
    -> Hash  (cost=1063.79..1063.79 rows=17779 width=13)
        (actual time=8.667..8.667 rows=17720 loops=1)
        Buckets: 32768  Batches: 1  Memory Usage: 1087kB
    -> Seq Scan on subjects s2
        (cost=0.00..1063.79 rows=17779 width=13)
        (actual time=0.009..4.677 rows=17779 loops=1)
Planning time: 0.255 ms
Execution time: 1191.023 ms

```

... EXPLAIN Examples

85/141

Example: Join on a non-indexed attribute (2018)

```

uni=# explain analyze
uni=# select s1.code, s2.code
uni=# from Subjects s1, Subjects s2
uni=# where s1.offeredBy = s2.offeredBy and s1.code < s2.code;
               QUERY PLAN
-----
Hash Join  (cost=1286.03..126135.12 rows=2371100 width=18)
    (actual time=7.356..6806.042 rows=3655437 loops=1)
    Hash Cond: (s1.offeredby = s2.offeredby)
    Join Filter: (s1.code < s2.code)
    Rows Removed by Join Filter: 3673157
    -> Seq Scan on subjects s1
        (cost=0.00..1063.79 rows=17779 width=13)
        (actual time=0.009..4.602 rows=17779 loops=1)
    -> Hash  (cost=1063.79..1063.79 rows=17779 width=13)
        (actual time=7.301..7.301 rows=17720 loops=1)
        Buckets: 32768  Batches: 1  Memory Usage: 1087kB
    -> Seq Scan on subjects s2
        (cost=0.00..1063.79 rows=17779 width=13)
        (actual time=0.005..4.452 rows=17779 loops=1)
Planning time: 0.159 ms
Execution time: 6949.167 ms

```

Exercise 4: EXPLAIN examples

86/141

Using the database described earlier ...

```

Course_enrolments(student, course, mark, grade, ...)
Courses(id, subject, semester, homepage)

```



```

People(id, family, given, title, name, ..., birthday)
Program_enrolments(id, student, semester, program, wam, ...)
Students(id, stype)
Subjects(id, code, name, longname, uoc, offeredby, ...)

```

```

create view EnrolmentCounts as
select s.code, c.semester, count(e.student) as nstudes
  from Courses c join Subjects s on c.subject=s.id
    join Course_enrolments e on e.course = c.id
 group by s.code, c.semester;

```

predict how each of the following queries will be executed ...

Check your prediction using the EXPLAIN ANALYZE command.

1. select max(birthday) from People
2. select max(id) from People
3. select family from People order by family
4. select distinct p.id, pname
from People s, CourseEnrolments e
where s.id=e.student and e.grade='FL'
5. select * from EnrolmentCounts where code='COMP9315';

Examine the effect of adding ORDER BY and DISTINCT.

Add indexes to improve the speed of slow queries.

Transaction Processing

Transaction Processing

89/141

Transaction (tx) = application-level atomic op, multiple DB ops

Concurrent transactions are

- desirable, for improved performance
- problematic, because of potential unwanted interactions

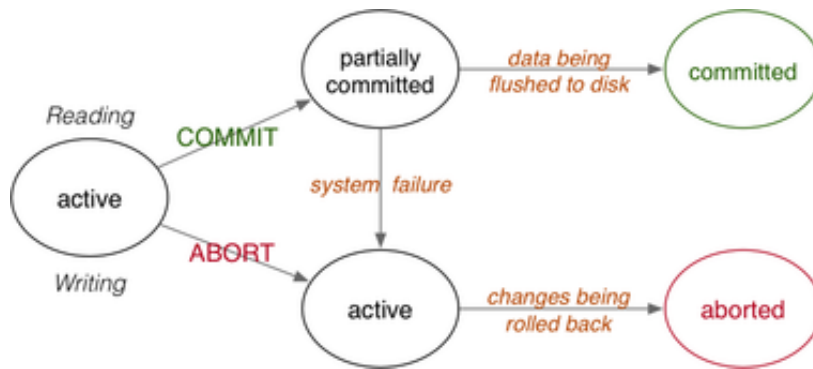
To ensure problem-free concurrent transactions:

- **Atomic** ... whole effect of tx, or nothing
 - **Consistent** ... individual tx's are "correct" (wrt application)
 - **Isolated** ... each tx behaves as if no concurrency
 - **Durable** ... effects of committed tx's persist
-

... Transaction Processing

90/141

Transaction states:



COMMIT \Rightarrow all changes preserved, ABORT \Rightarrow database unchanged

... Transaction Processing

91/141

Transaction processing:

- the study of techniques for realising ACID properties

Consistency is the property:

- a tx is correct with respect to its own specification
- a tx performs a mapping that maintains all DB constraints

Ensuring this must be left to application programmers.

Our discussion focusses on: Atomicity, Durability, Isolation

... Transaction Processing

92/141

Atomicity is handled by the *commit* and *abort* mechanisms

- **commit** ends tx and ensures all changes are saved
- **abort** ends tx and *undoes* changes "already made"

Durability is handled by implementing *stable storage*, via

- redundancy, to deal with hardware failures
- logging/checkpoint mechanisms, to recover state

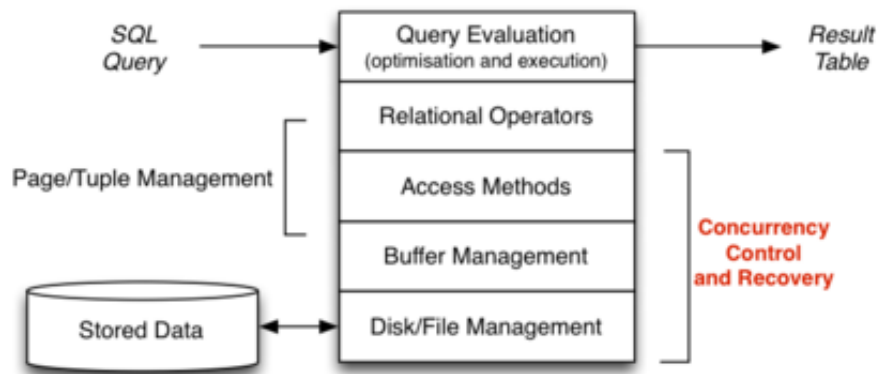
Isolation is handled by *concurrency control* mechanisms

- two possibilities: lock-based, timestamp-based
- various levels of isolation are possible (e.g. serializable)

... Transaction Processing

93/141

Where transaction processing fits in the DBMS:



Transaction Terminology

94/141

To describe transaction effects, we consider:

- **READ** - transfer data from "disk" to memory
- **WRITE** - transfer data from memory to "disk"
- **ABORT** - terminate transaction, unsuccessfully
- **COMMIT** - terminate transaction, successfully

Relationship between the above operations and SQL:

- **SELECT** produces **READ** operations on the database
- **UPDATE** and **DELETE** produce **READ** then **WRITE** operations
- **INSERT** produces **WRITE** operations

... Transaction Terminology

95/141

More on transactions and SQL

- **BEGIN** starts a transaction
 - the `begin` keyword in PLpgSQL is not the same thing
- **COMMIT** commits and ends the current transaction
 - some DBMSs e.g. PostgreSQL also provide `END` as a synonym
 - the `end` keyword in PLpgSQL is not the same thing
- **ROLLBACK** aborts the current transaction, undoing any changes
 - some DBMSs e.g. PostgreSQL also provide `ABORT` as a synonym

In PostgreSQL, tx's cannot be defined inside stored procedures (e.g. PLpgSQL)

... Transaction Terminology

96/141

The **READ**, **WRITE**, **ABORT**, **COMMIT** operations:

- occur in the context of some transaction T
- involve manipulation of data items X, Y, \dots (**READ** and **WRITE**)

The operations are typically denoted as:

$R_T(X)$ read item X in transaction T

$W_T(X)$ write item X in transaction T

A_T abort transaction T

C_T commit transaction T

Schedules

97/141

A *schedule* gives the sequence of operations from ≥ 1 tx

Serial schedule for a set of tx's $T_1 \dots T_n$

- all operations of T_i complete before T_{i+1} begins

E.g. $R_{T_1}(A)$ $W_{T_1}(A)$ $R_{T_2}(B)$ $R_{T_2}(A)$ $W_{T_3}(C)$ $W_{T_3}(B)$

Concurrent schedule for a set of tx's $T_1 \dots T_n$

- operations from individual T_i 's are interleaved

E.g. $R_{T_1}(A)$ $R_{T_2}(B)$ $W_{T_1}(A)$ $W_{T_3}(C)$ $W_{T_3}(B)$ $R_{T_2}(A)$

... Schedules

98/141

Serial schedules guarantee database consistency

- each T_i commits before T_{i+1}
- prior to T_i database is consistent
- after T_i database is consistent (assuming T_i is correct)
- before T_{i+1} database is consistent ...

Concurrent schedules interleave operations arbitrarily

- and may produce a database that is not consistent
- after all of the transactions have committed

Transaction Anomalies

99/141

What problems can occur with uncontrolled concurrent transactions?

The set of phenomena can be characterised broadly under:

- dirty read*:
reading data item currently in use by another tx
- nonrepeatable read*:
re-reading data item, since changed by another tx
- phantom read*:
re-reading result set, since changed by another tx

... Transaction Anomalies

100/141

Dirty read: a transaction reads data written by a concurrent uncommitted transaction

Example:

Transaction T1	Transaction T2
(1) select a into X	

```

        from R where id=1
(2)          select a into Y
              from R where id=1
(3)  update R set a=X+1
      where id=1
(4)  commit
(5)          update R set a=Y+1
              where id=1
(6)  commit

```

Effect: T1's update on R.a is lost.

... Transaction Anomalies

101/141

Nonrepeatable read: a transaction re-reads data it has previously read and finds that data has been modified by another transaction (that committed since the initial read)

Example:

	Transaction T1	Transaction T2
(1)	select * from R where id=5	
(2)		update R set a=8 where id=5
(3)		commit
(4)	select * from R where id=5	

Effect: T1 runs same query twice; sees different data

... Transaction Anomalies

102/141

Phantom read: a transaction re-executes a query returning a set of rows that satisfy a search condition and finds that the set of rows satisfying the condition has changed due to another recently-committed transaction

Example:

	Transaction T1	Transaction T2
(1)	select count(*) from R where a=5	
(2)		insert into R(id,a,b) values (2,5,8)
(3)		commit
(4)	select count(*) from R where a=5	

Effect: T1 runs same query twice; sees different result set

Example of Transaction Failure

103/141

Above examples assumed that all transactions commit.

Additional problems can arise when transactions abort.

Consider the following schedule where transaction T1 fails:

T1: R(X) W(X) A
T2: R(X) W(X) C

Abort *will* rollback the changes to x, but ...

Consider three places where rollback might occur:

T1: R(X) W(X) A [1] [2] [3]
 T2: R(X) W(X) C

... Example of Transaction Failure

104/141

Abort / rollback scenarios:

T1: R(X) W(X) A [1] [2] [3]
 T2: R(X) W(X) C

Case [1] is ok

- all effects of T1 vanish; final effect is simply from T2

Case [2] is problematic

- some of T1's effects persist, even though T1 aborted

Case [3] is also problematic

- T2's effects are lost, even though T2 committed

Transaction Isolation


106/141

Simplest form of isolation: *serial* execution ($T_1; T_2; T_3; \dots$)

Problem: serial execution yields poor throughput.

Concurrency control schemes (CCSs) aim for "safe" concurrency

Abstract view of DBMS concurrency mechanisms:



Serializability

107/141

Consider two schedules S_1 and S_2 produced by

- executing the same set of transactions $T_1..T_n$ concurrently
- but with a non-serial interleaving of *R/W* operations

S_1 and S_2 are *equivalent* if $StateAfter(S_1) = StateAfter(S_2)$

- i.e. final state yielded by S_1 is same as final state yielded by S_2

S is a *serializable schedule* (for a set of concurrent tx's $T_1 .. T_n$) if

- S is equivalent to some serial schedule S_s of $T_1 .. T_n$

Under these circumstances, consistency is guaranteed
(assuming no aborted transactions and no system failures)

... Serializability

108/141

Two formulations of serializability:

- *conflict serializability*
 - i.e. conflicting R/W operations occur in the "right order"
 - check via precedence graph; look for absence of cycles
- *view serializability*
 - i.e. read operations see the correct version of data
 - checked via VS conditions on likely equivalent schedules

View serializability is strictly weaker than conflict serializability.

Exercise 5: Serializability Checking

109/141

Is the following schedule view/conflict serializable?

```

T1:      W(B)  W(A)
T2:  R(B)              W(A)
T3:              R(A)      W(A)

```

Is the following schedule view/conflict serializable?

```

T1:      W(B)  W(A)
T2:  R(B)              W(A)
T3:              R(A)  W(A)

```

Transaction Isolation Levels

110/141

SQL programmers' concurrency control mechanism ...

```

set transaction
  read only -- so weaker isolation may be ok
  read write -- suggests stronger isolation needed
isolation level
  -- weakest isolation, maximum concurrency
  read uncommitted
  read committed
  repeatable read
  serializable
  -- strongest isolation, minimum concurrency

```

Applies to current tx only; affects how scheduler treats this tx.

... Transaction Isolation Levels

111/141

Implication of transaction isolation levels:

Isolation	Dirty	Nonrepeatable	Phantom
-----------	-------	---------------	---------

Level	Read	Read	Read
Read Uncommitted	Possible	Possible	Possible
Read Committed	Not Possible	Possible	Possible
Repeatable Read	Not Possible	Not Possible	Possible
Serializable	Not Possible	Not Possible	Not Possible

... Transaction Isolation Levels

112/141

For transaction isolation, PostgreSQL

- provides syntax for all four levels
- treats *read uncommitted* as *read committed*
- *repeatable read* behaves like *serializable*
- default level is *read committed*

Note: cannot implement *read uncommitted* because of MVCC

For more details, see [PostgreSQL Documentation section 13.2](#)

- extensive discussion of semantics of UPDATE, INSERT, DELETE

... Transaction Isolation Levels

113/141

A PostgreSQL tx consists of a sequence of SQL statements:

```
BEGIN  $S_1$ ;  $S_2$ ; ...  $S_n$ ; COMMIT;
```

Isolation levels affect view of DB provided to each S_i :

- in *read committed* ...
 - each S_i sees snapshot of DB at start of S_i
- in *repeatable read* and *serializable* ...
 - each S_i sees snapshot of DB at start of tx
 - serializable checks for extra conditions

Transactions fail if the system detects violation of isolation level.

... Transaction Isolation Levels

114/141

Example of *repeatable read* vs *serializable*

- table R(class,value) containing (1,10) (1,20) (2,100) (2,200)
- **T1**: X = sum(value) where class=1; insert R(2,X); commit
- **T2**: X = sum(value) where class=2; insert R(1,X); commit
- with *repeatable read*, both transactions commit, giving
 - updated table: (1,10) (1,20) (2,100) (2,200) (1,300) (2,30)
- with *serial* transactions, only one transaction commits

- T1;T2 gives (1,10) (1,20) (2,100) (2,200) (2,30) (1,330)
- T2;T1 gives (1,10) (1,20) (2,100) (2,200) (1,300) (2,330)
- PG recognises that committing both gives serialization violation

Implementing Concurrency Control

Concurrency Control

116/141

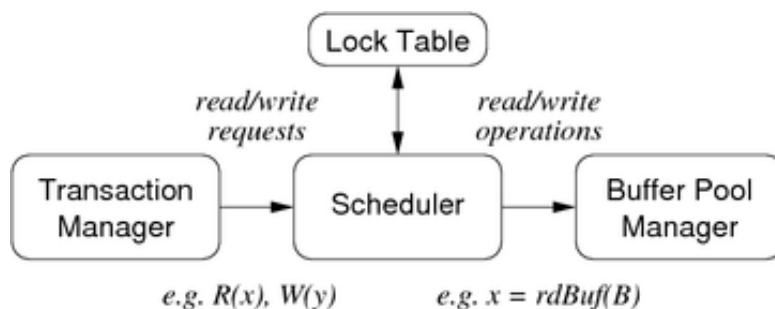
Approaches to concurrency control:

- *Lock-based*
 - Synchronise tx execution via locks on relevant part of DB.
- *Version-based* (multi-version concurrency control)
 - Allow multiple consistent versions of the data to exist.
 - Each tx has access only to version existing at start of tx.
- *Validation-based* (optimistic concurrency control)
 - Execute all tx's; check for validity problems on commit.
- *Timestamp-based*
 - Organise tx execution via timestamps assigned to actions.

Lock-based Concurrency Control

117/141

Locks introduce additional mechanisms in DBMS:



The Lock Manager

- manages the locks requested by the scheduler

... Lock-based Concurrency Control

118/141

Lock table entries contain:

- object being locked (DB, table, tuple, field)
- type of lock: read/shared, write/exclusive
- FIFO queue of tx's requesting this lock
- count of tx's currently holding lock (max 1 for write locks)

Lock and unlock operations *must* be atomic.

Lock *upgrade*:

- if a tx holds a read lock, and it is the only tx holding that lock
- then the lock can be converted into a write lock

... Lock-based Concurrency Control

119/141

Synchronise access to shared data items via following rules:

- before reading X , get read (shared) lock on X
- before writing X , get write (exclusive) lock on X
- a tx attempting to get a read lock on X is blocked if another tx already has write lock on X
- a tx attempting to get an write lock on X is blocked if another tx has any kind of lock on X

These rules alone do not guarantee serializability.

... Lock-based Concurrency Control

120/141

Consider the following schedule, using locks:

T1(a): $L_r(Y)$ $R(Y)$ continued

T2(a): $L_r(X)$ $R(X)$ $U(X)$ continued

T1(b): $U(Y)$ $L_w(X)$ $W(X)$ $U(X)$

T2(b): $L_w(Y)$ $W(Y)$ $U(Y)$

(where L_r = read-lock, L_w = write-lock, U = unlock)

Locks correctly ensure controlled access to x and y .

Despite this, the schedule is not serializable. (Ex: prove this)

Two-Phase Locking

121/141

To guarantee serializability, we require an additional constraint:

- in every tx, all *lock* requests precede all *unlock* requests

Each transaction is then structured as:

- *growing* phase where locks are acquired
- *action* phase where "real work" is done
- *shrinking* phase where locks are released

Clearly, this reduces potential concurrency ...

Problems with Locking

122/141

Appropriate locking can guarantee correctness.

However, it also introduces potential undesirable effects:

- *Deadlock*
 - No transactions can proceed; each waiting on lock held by another.
- *Starvation*
 - One transaction is permanently "frozen out" of access to data.
- *Reduced performance*
 - Locking introduces delays while waiting for locks to be released.

Deadlock

123/141

Deadlock occurs when two transactions are waiting for a lock on an item held by the other.

Example:

T1: $L_w(A)$ $R(A)$ $L_w(B)$
 T2: $L_w(B)$ $R(B)$ $L_w(A)$

How to deal with deadlock?

- prevent it happening in the first place
- let it happen, detect it, recover from it

... Deadlock

124/141

Handling deadlock involves forcing a transaction to "back off".

- select process to "back off"
 - choose on basis of how far transaction has progressed, # locks held, ...
- roll back the selected process
 - how far does this it need to be rolled back? (less roll-back is better)
 - worst-case scenario: abort one transaction
- prevent starvation
 - need methods to ensure that same transaction isn't always chosen

... Deadlock

125/141

Methods for managing deadlock

- *timeout* : set max time limit for each tx
- *waits-for graph* : records T_j waiting on lock held by T_k
 - *prevent* deadlock by checking for new cycle \Rightarrow abort T_i
 - *detect* deadlock by periodic check for cycles \Rightarrow abort T_i
- *timestamps* : use tx start times as basis for priority
 - scenario: T_j tries to get lock held by T_k ...
 - *wait-die*: if $T_j < T_k$, then T_j waits, else T_j rolls back
 - *wound-wait*: if $T_j < T_k$, then T_k rolls back, else T_j waits

... Deadlock

126/141

Properties of deadlock handling methods:

- both wait-die and wound-wait are fair
- wait-die tends to
 - roll back tx's that have done little work
 - but rolls back tx's more often
- wound-wait tends to
 - roll back tx's that may have done significant work
 - but rolls back tx's less often
- timestamps easier to implement than waits-for graph
- waits-for minimises roll backs because of deadlock

Exercise 6: Deadlock Handling

127/141

Consider the following schedule on four transactions:

T1: $R(A)$ $W(C)$ $W(D)$
 T2: $R(B)$ $W(C)$
 T3: $R(D)$ $W(B)$

T4 : R (E) W (A)

Assume that: each R acquires a shared lock; each W uses an exclusive lock; two-phase locking is used.

Show how the wait-for graph for the locks evolves.

Show how any deadlocks might be resolved via this graph.

Optimistic Concurrency Control

128/141

Locking is a pessimistic approach to concurrency control:

- limit concurrency to ensure that conflicts don't occur

Costs: lock management, deadlock handling, contention.

In scenarios where there are far more reads than writes ...

- don't lock (allow arbitrary interleaving of operations)
- check just before commit that no conflicts occurred
- if problems, roll back conflicting transactions


... Optimistic Concurrency Control

129/141

Transactions have three distinct phases:

- *Reading*: read from database, modify local copies of data
- *Validation*: check for conflicts in updates
- *Writing*: commit local copies of data to database

Timestamps are recorded at points noted:



... Optimistic Concurrency Control

130/141

Data structures needed for validation:

- A ... set of txs that are reading data and computing results
- V ... set of txs that have reached validation (not yet committed)
- F ... set of txs that have finished (committed data to storage)
- for each T_i , timestamps for when it reached A , V , F
- $R(T_i)$ set of all data items read by T_i
- $W(T_i)$ set of all data items to be written by T_i

Use the V timestamps as ordering for transactions

- assume serial tx order based on ordering of $V(T_i)$'s

... Optimistic Concurrency Control

131/141

Validation check for transaction T

- for all transactions $T_i \neq T$
 - if $V(T_i) < A(T) < F(T_i)$, then check $W(T_i) \cap R(T)$ is empty
 - if $V(T_i) < V(T) < F(T_i)$, then check $W(T_i) \cap W(T)$ is empty

If this check fails for any T_i , then T is rolled back.

Prevents: T reading dirty data, T overwriting T_i 's changes

Problems: rolls back "complete" tx's, cost to maintain A, V, F sets

Multi-version Concurrency Control

132/141

Multi-version concurrency control (MVCC) aims to

- retain benefits of locking, while getting more concurrency
- by providing multiple (consistent) versions of data items

Achieves this by

- readers access an "appropriate" version of each data item
- writers make new versions of the data items they modify

Main difference between MVCC and standard locking:

- read locks do not conflict with write locks \Rightarrow
- reading never blocks writing, writing never blocks reading

... Multi-version Concurrency Control

133/141

WTS = timestamp of last writer; RTS = timestamp of last reader

Chained tuple versions: $tup_{oldest} \rightarrow tup_{older} \rightarrow tup_{newest}$

When a reader T_i is accessing the database

- ignore any data item D created after T_i started ($WTS(D) > TS(T_i)$)
- use only newest version V satisfying $WTS(V) < TS(T_i)$

When a writer T_j attempts to change a data item

- find newest version V satisfying $WTS(V) < TS(T_j)$
- if no later versions exist, create new version of data item
- otherwise, abort T_j

... Multi-version Concurrency Control

134/141

Advantage of MVCC

- locking needed for serializability considerably reduced

Disadvantages of MVCC

- visibility-check overhead (on every tuple read/write)
- reading an item V causes an update of $RTS(V)$
- storage overhead for extra versions of data items
- overhead in removing out-of-date versions of data items

Despite apparent disadvantages, MVCC is very effective.

... Multi-version Concurrency Control

135/141

Removing old versions:

- V_j and V_k are versions of same item
- $WTS(V_j)$ and $WTS(V_k)$ precede $TS(T_i)$ for all T_i
- remove version with smaller $WTS(V_x)$ value

When to make this check?

- every time a new version of a data item is added?
- periodically, with fast access to blocks of data

PostgreSQL uses the latter (*vacuum*).

Concurrency Control in PostgreSQL

136/141

PostgreSQL uses two styles of concurrency control:

- multi-version concurrency control (MVCC)
(used in implementing SQL DML statements (e.g. `select`))
- two-phase locking (2PL)
(used in implementing SQL DDL statements (e.g. `create table`))

From the SQL (PLpgSQL) level:

- can let the lock/MVCC system handle concurrency
- can handle it explicitly via `LOCK` statements

... Concurrency Control in PostgreSQL

137/141

PostgreSQL provides *read committed* and *serializable* isolation levels.

Using the serializable isolation level, a `select`:

- sees only data committed before the transaction began
- never sees changes made by concurrent transactions

Using the serializable isolation level, an update fails:

- if it tries to modify an "active" data item
(active = affected by some other tx, either committed or uncommitted)

The transaction containing the update must then rollback and re-start.

... Concurrency Control in PostgreSQL

138/141

Implementing MVCC in PostgreSQL requires:

- a log file to maintain current status of each T_i
- in every tuple:
 - `xmin` ID of the tx that created the tuple
 - `xmax` ID of the tx that replaced/deleted the tuple (if any)
 - `xnew` link to newer versions of tuple (if any)

- for each transaction T_i :
 - a transaction ID (timestamp)
 - SnapshotData: list of active tx's when T_i started

... Concurrency Control in PostgreSQL

139/141

Rules for a tuple to be visible to T_i :

- the `xmin` (creation transaction) value must
 - be committed in the log file
 - have started before T_i 's start time
 - not be active at T_i 's start time
- the `xmax` (delete/replace transaction) value must
 - be blank or refer to an aborted tx, or
 - have started after T_i 's start time, or
 - have been active at SnapshotData time

For details, see: [utils/time/tqual.c](#)

... Concurrency Control in PostgreSQL

140/141

Tx's always see a consistent version of the database.

But may not see the "current" version of the database.

E.g. T_1 does select, then concurrent T_2 deletes some of T_1 's selected tuples

This is OK unless tx's communicate outside the database system.

E.g. T_1 counts tuples while T_2 deletes then counts; then counts are compared

Use locks if application needs every tx to see same current version

- `LOCK TABLE` locks an entire table
- `SELECT FOR UPDATE` locks only the selected rows

Exercise 7: Locking in PostgreSQL

141/141

How could we solve this problem via locking?

```
create or replace function
  allocSeat(paxID int, fltID int, seat text)
  returns boolean
as $$
declare
  pid int;
begin
  select paxID into pid from SeatingAlloc
  where flightID = fltID and seatNum = seat;
  if (pid is not null) then
    return false; -- someone else already has seat
  else
    update SeatingAlloc set pax = paxID
    where flightID = fltID and seatNum = seat;
    commit;
    return true;
  end if;
end;
```

```
$$ language plpgsql;
```

Produced: 13 Sep 2018